



Evaluation Metrics Documentation

Hybrid RAG System with Automated Evaluation

GitHub Repository: https://github.com/vishalvishal099/Hybrid_RAG_System_with_Automated_Evaluation

Table of Contents

1. Mandatory Metric: MRR (Mean Reciprocal Rank)
 2. Custom Metric 1: Recall@10
 3. Custom Metric 2: Answer F1 Score
 4. Metric Comparison Summary
 5. Implementation References
-

1. Mandatory Metric: MRR (Mean Reciprocal Rank)

1.1 Definition

Mean Reciprocal Rank (MRR) measures how well a retrieval system ranks the first relevant document. It is the average of reciprocal ranks across all queries.

1.2 Mathematical Formulation

$$\text{MRR} = \frac{1}{|Q|} \sum_{i=1}^{|Q|} \frac{1}{\text{rank}_i}$$

Where:

- $|Q|$ = Total number of queries
- rank_i = Rank position of the first relevant document for query i
- If no relevant document is found, $\frac{1}{\text{rank}_i} = 0$

1.3 Example Calculation

Query	Relevant URL Rank	Reciprocal Rank
Q1	1	$1/1 = 1.000$
Q2	3	$1/3 = 0.333$
Q3	2	$1/2 = 0.500$
Q4	Not Found	0.000

$$\text{MRR} = \frac{1.0 + 0.333 + 0.5 + 0.0}{4} = 0.458$$

1.4 Why MRR for RAG?

1. **Focus on Top Result:** In RAG systems, the first relevant chunk heavily influences answer quality

2. **User Experience:** Users typically examine top results first
3. **Penalizes Poor Ranking:** Low ranks receive exponentially smaller scores

1.5 Interpretation Guidelines

MRR Score	Interpretation	System Quality
0.90 - 1.00	Excellent	Relevant docs almost always rank first
0.70 - 0.89	Good	Relevant docs typically in top 2
0.50 - 0.69	Moderate	Relevant docs usually in top 3-4
0.30 - 0.49	Fair	Relevant docs often below top 3
0.00 - 0.29	Poor	Significant ranking issues

1.6 Our Results

Method	MRR Score	Interpretation
Dense	0.3025	Fair - relevant docs often below top 3
Sparse (BM25)	0.4392	Fair-Moderate - best performance
Hybrid (RRF)	0.3783	Fair - between dense and sparse

2. Custom Metric 1: Recall@10

2.1 Justification for Selection

Why Recall@10 was chosen:

1. **Relevance to RAG Context Windows:** Our system retrieves top-10 chunks for answer generation. Recall@10 directly measures how many relevant documents appear in this context window.
2. **Complementary to MRR:** While MRR focuses on rank of first relevant result, Recall@10 measures coverage of all relevant documents.
3. **Answer Quality Correlation:** Higher Recall@10 means more relevant information available to the LLM for answer generation.
4. **Standard in IR Research:** Widely used metric enabling comparison with other systems.
5. **Practical Threshold:** 10 documents represent a reasonable context size for LLM processing without token overflow.

2.2 Mathematical Formulation

$$\text{Recall@K} = \frac{\text{|Retrieved_K |cap Relevant|}}{\text{|Relevant|}}$$

For our implementation:

$\text{Recall@10} = \frac{\text{Number of relevant URLs in top 10 results}}{\text{Total number of relevant URLs for query}}$

2.3 Detailed Calculation

Step 1: Retrieve top 10 chunks for query **Step 2:** Extract unique source URLs from retrieved chunks **Step 3:** Compare with ground truth relevant URLs **Step 4:** Calculate intersection ratio

Example:

```
Ground Truth URLs: [url_A, url_B]
Retrieved URLs (top 10): [url_A, url_C, url_D, url_E, ...]

Intersection: {url_A}
Recall@10 = 1/2 = 0.5 (50%)
```

2.4 Implementation

Code Reference: [evaluate_chromadb_fast.py](#)

```
def compute_recall_at_k(retrieved_urls, relevant_urls, k=10):
    """
    Compute Recall@K

    Args:
        retrieved_urls: List of URLs from top K retrieved chunks
        relevant_urls: List of ground truth relevant URLs
        k: Number of top results to consider (default: 10)

    Returns:
        float: Recall score between 0 and 1
    """
    if not relevant_urls:
        return 0.0

    retrieved_set = set(retrieved_urls[:k])
    relevant_set = set(relevant_urls)

    intersection = retrieved_set & relevant_set
    recall = len(intersection) / len(relevant_set)

    return recall
```

2.5 Interpretation Guidelines

Recall@10 Score	Interpretation	Implication
0.90 - 1.00	Excellent	Almost all relevant docs in top 10

Recall@10 Score	Interpretation	Implication
0.70 - 0.89	Good	Most relevant docs retrieved
0.50 - 0.69	Moderate	About half of relevant docs found
0.30 - 0.49	Fair	Significant relevant docs missing
0.00 - 0.29	Poor	Most relevant docs not in top 10

2.6 Our Results

Method	Recall@10	Interpretation
Dense	0.33	Fair - 1/3 of relevant docs retrieved
Sparse (BM25)	0.47	Fair-Moderate - best coverage
Hybrid (RRF)	0.43	Fair - improved over dense

3. Custom Metric 2: Answer F1 Score

3.1 Justification for Selection

Why Answer F1 was chosen:

1. **End-to-End Quality Assessment:** Unlike retrieval metrics, F1 evaluates the final generated answer quality.
2. **Token-Level Precision and Recall:** Captures both what the system correctly included (precision) and what relevant information was covered (recall).
3. **Handles Partial Matches:** Unlike exact match, F1 gives credit for partially correct answers.
4. **Standard NLP Metric:** Widely used in QA systems (SQuAD, etc.), enabling benchmarking.
5. **Balances Conciseness and Completeness:** Precision prevents over-generation; recall ensures coverage.

3.2 Mathematical Formulation

$$\text{Precision} = \frac{|\text{Generated} \cap \text{Reference}|}{|\text{Generated}|}$$

$$\text{Recall} = \frac{|\text{Generated} \cap \text{Reference}|}{|\text{Reference}|}$$

$$F1 = 2 \times \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}$$

Where:

- Generated = Set of tokens in generated answer (after normalization)
- Reference = Set of tokens in ground truth answer

3.3 Detailed Calculation

Preprocessing Steps:

1. Convert to lowercase
2. Remove punctuation
3. Tokenize into words
4. Remove stopwords (optional)

Example:

Generated Answer: "The capital of France is Paris"
 Reference Answer: "Paris is the capital city of France"

Generated Tokens: {the, capital, of, france, is, paris}
 Reference Tokens: {paris, is, the, capital, city, of, france}

Common Tokens: {the, capital, of, france, is, paris} = 6
 Generated Count: 6
 Reference Count: 7

Precision = $6/6 = 1.00$
 Recall = $6/7 = 0.857$
 $F1 = 2 \times (1.00 \times 0.857) / (1.00 + 0.857) = 0.923$

3.4 Implementation

Code Reference: [evaluate_chromadb_fast.py](#)

```
def compute_answer_f1(generated_answer: str, reference_answer: str) ->
    float:
    """
    Compute token-level F1 score between generated and reference answers.

    Args:
        generated_answer: Model's generated answer
        reference_answer: Ground truth answer

    Returns:
        float: F1 score between 0 and 1
    """
    def normalize(text):
        """Normalize text: lowercase, remove punctuation, tokenize"""
        text = text.lower()
        text = re.sub(r'[^w\s]', '', text)
        tokens = text.split()
        return set(tokens)

    gen_tokens = normalize(generated_answer)
    ref_tokens = normalize(reference_answer)

    if not gen_tokens or not ref_tokens:
```

```

    return 0.0

common = gen_tokens & ref_tokens

precision = len(common) / len(gen_tokens) if gen_tokens else 0
recall = len(common) / len(ref_tokens) if ref_tokens else 0

if precision + recall == 0:
    return 0.0

f1 = 2 * (precision * recall) / (precision + recall)
return f1

```

3.5 Interpretation Guidelines

Answer F1 Score	Interpretation	Answer Quality
0.80 - 1.00	Excellent	Nearly identical to reference
0.60 - 0.79	Good	Most key information present
0.40 - 0.59	Moderate	Partial overlap with reference
0.20 - 0.39	Fair	Some relevant tokens, major gaps
0.00 - 0.19	Poor	Minimal overlap with reference

3.6 Our Results

Method	Avg Answer F1	Interpretation
Dense	~0.05	Poor - answers don't match reference
Sparse (BM25)	~0.05	Poor - answers don't match reference
Hybrid (RRF)	~0.05	Poor - answers don't match reference

Analysis of Low F1 Scores: The low F1 scores indicate a mismatch between generated answers and reference answers. This is often due to:

1. **Answer Style Differences:** LLM generates verbose answers vs. concise references
2. **Paraphrasing:** Same meaning but different words
3. **Context Window Issues:** Retrieved context doesn't contain answer information

4. Metric Comparison Summary

4.1 Metric Properties

Property	MRR	Recall@10	Answer F1
Measures	Ranking Quality	Coverage	Answer Quality

Property	MRR	Recall@10	Answer F1
Focus	First relevant result	All relevant results	Generated text
Range	0 to 1	0 to 1	0 to 1
Higher is	Better	Better	Better
Stage	Retrieval	Retrieval	Generation

4.2 Combined Insights

Method	MRR	Recall@10	Answer F1	Overall
Dense	0.30	0.33	0.05	Weak retrieval, poor answers
Sparse	0.44	0.47	0.05	Best retrieval , poor answers
Hybrid	0.38	0.43	0.05	Moderate retrieval, poor answers

4.3 Key Findings

- BM25 (Sparse) dominates:** Outperforms dense embeddings on Wikipedia-style factual content
- Hybrid provides balance:** RRF fusion gives intermediate performance
- Answer generation is the bottleneck:** Low F1 scores across all methods indicate LLM generation needs improvement
- Retrieval ≠ Answer Quality:** Good retrieval (Recall@10 = 0.47) doesn't guarantee good answers (F1 = 0.05)

5. Implementation References

5.1 Code Files

Component	File	GitHub Link
Evaluation Script	evaluate_chromadb_fast.py	View
RAG System	chromadb_rag_system.py	View
Metrics Library	evaluation/comprehensive_metrics.py	View

5.2 Result Files

File	Purpose	GitHub Link
evaluation_results_chromadb.csv	Detailed results (300 rows)	View
evaluation_summary_chromadb.json	Summary statistics	View